

I. Introduction to Reinforcement Learning

1. Introduction

Reinforcement learning (RL) is a science of *decision-making*. We are going to discuss RL from the perspective of machine learning (ML).

What makes RL different from other ML algorithms (e.g. supervised/unsupervised learning)? [1.1]

- (a) There is no supervisor, only a *reward* signal.
- (b) Feedback is delayed, not instantaneous.
- (c) Time really matters (i.e. sequential, non iid data is supplied).
- (d) *Agent's* action affect the subsequent data it receives.

Example 1.1. Examples of RL

- (a) Fly stunt manoeuvring in a helicopter.
 - (b) Playing Backgammon.
 - (c) Managing an investment portfolio.
 - (d) Controlling a power station.
 - (e) Making a humanoid robot walk.
 - (f) Playing Atari games.
-

2. The RL Problem

Def'n 1.1. **Reward**

A **reward** R_t at **step** t is a scalar feedback signal indicating how well agent is doing at step t .

The job of an RL agent is to maximize cumulative reward. RL is based on the following hypothesis.

Statement 1.1. Reward Hypothesis

Goals can be described by maximization of expected cumulative reward.

As a consequence of the above hypothesis, we have:

- (a) Actions may have long term sequence.
- (b) Reward may be delayed. *It may be better to sacrifice immediate reward to gain more long-term reward.*

For instance, in a financial investment, it may take months to mature.

We can describe RL by using the relationship between agent and environment: at each step t , the agent

- (a) executes action A_t (towards the environment); and
- (b) receives observation O_t and scalar reward R_t from the environment.

We repeat this over and over to train our agent.

Def'n 1.2. **History**

The **history** H_t at step t is the sequence of observations, actions, rewards up to time t :

$$H_t = ((A_i, O_i, R_i))_{i=1}^t = (A_1, O_1, R_1), \dots, (A_t, O_t, R_t).$$

What happens next depends on the history. The agent select actions and the environment selects observations and rewards. But history is often too long, so instead we use less amount of information to determine what happens next.

Def'n 1.3. **State**

The **state** S_t at step t is the information used to determine what happens at step $t + 1$.

There are three parts of state.

Def'n 1.4. **Environmental State**

The **environmental state** S_t^e at step t is the environment's private representation that spits out next observation and reward.

The environment state is not usuall visible to the agent. Even if it is visible, it may contain irrelevant information.

Def'n 1.5. **Agent State**

The **agent state** S_t^a is the agent's internal representation the agent uses to pick the next action. Formally, it is a function of the history:

$$S_t = f(H_t)$$

for some function f .

When we speak of state, we shall always mean agent state.

Def'n 1.6. **Markov** Sequence of State

We say the sequence of states $(S_t)_{t=1}$ is **Markov** if

$$\mathbb{P}(S_{t+1}|S_1, \dots, S_t) = \mathbb{P}(S_{t+1}|S_t)$$

for all t .

In words, for a Markov process,

$$\text{the future independent of the past given the present.} \quad [1.2]$$

In other words, once S_t is known, the history may be thrown away (i.e. the state S_t is a sufficient statistic of the future). By definition, the environment state S_t^e and the history H_t are Markov.

Def'n 1.7. **Fully Observable, Partially Observable** Environment

A **fully observable** environment is when the agent directly observes the environmental stae. That is,

$$S_t^a = S_t^e.$$

Otherwise, we say that the environment is **partially observable**.

In a fully observable environment, $(S_t)_{t=1}$ is a *Markov decision process* (MDP), to which we will come back shortly.

3. Inside an RL Agent

An RL agent may include one or more of these components.

Def'n 1.8. **Policy, Value Function, Model** of an RL Agent

The **policy** π of an RL agent is the agent's behaviour function. A **deterministic** policy is a deterministic function from state s to action a :

$$a = \pi(s)$$

for all state s . A **stochastic policy** has the form of

$$\pi(a|s) = \mathbb{P}(A = a|S = s).$$

We desire π such that we can get back as much rewards back as possible from a .

The **value function** is the measure of how good is each state or action (i.e. a prediction of future reward). Formally,

$$v_{\pi}(s) = \mathbb{E}_{\pi}(R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots | S_t = s)$$

for some $\gamma \in [0, 1]$.¹

The **model** is the agent's representation of the environment that predicts what the environment will do next. There are two parts of a model. A **transition** model \mathcal{P} predicts the next state:

$$\mathcal{P}_{ss'}^a = \mathbb{P}(S' = s' | S = s, A = a).$$

A **reward** model \mathcal{R} predicts the next reward:

$$\mathcal{R}_s^a = \mathbb{E}(R | S = s, A = a).$$

¹The parameter γ says that we care about immediate rewards more than future rewards (unless future rewards are much larger).

Def'n 1.9. **Value-based, Policy-based, Actor Critic** RL Agent

We say an RL agent is

- (a) **value-based** if it has a value function but no policy;¹
- (b) **policy-based** if it has a policy but no value function; and
- (c) **actor critic** if it has both.

¹A value-based RL agent *reads off* value function *greedily* to come up with an implicit policy.

Def'n 1.10. **Model-free, Model-based** RL Agent

We say an RL agent is

- (a) **model-free** if it has no model; and
- (b) **model-based** if it has one.

4. Problems within RL

There are two fundamental problems in sequential decision making.

	RL	Planning
Environment	initially unknown	a model is known
Agent	interacts with the environment to improve its policy	performs computation with its model without external interactions to improve policy

In short,

RL is like a trial-and-error learning. [1.3]

The agent should discover a good policy from its experiences of the environment, without losing too much reward along the way. In doing so, *exploration* finds more information about the environment and *exploitation* exploits known information to maximize reward. We have to balance exploration and exploitation to get the best result.

In RL, we would like to *optimize the future* (i.e. find the best policy). This is called *control*. It turns out we have to make *predictions* (i.e. given a policy, evaluate the future) to do so.

II. Markov Decision Processes

As mentioned in Section 1, *Markov decision processes* (MDPs) formally describe an environment for RL when the environment is *fully observable*. That is, the current state completely characterizes the process. The nice thing about MDPs is that most RL problems can be formalized in to MDPs. For instance,

- (a) optimal control primarily deals with continuous MDPs;
- (b) partially observable problems can be converted in to MDPs; and
- (c) bandits are MDPs with one state.

1. Markov Processes

Let us start by recalling the definition of *Markov property*.

Recall 2.1. **Markov Process**

A **stochastic process** is a sequence $(X_t)_t$ of random variables.

We say a stochastic process $(S_t)_t$ is **Markov** if

$$\mathbb{P}(S_{t+1}|S_t, \dots, S_1) = \mathbb{P}(S_{t+1}|S_t).$$

For a Markov state s and successor state s' , the **state transition probability**, denoted as $\mathcal{P}_{ss'}$, is defined by

$$\mathcal{P}_{ss'} = \mathbb{P}(S_{t+1} = s' | S_t = s).$$

The **state transition matrix** \mathcal{P} defines transition probabilities from all states s to all successor states s' ,

$$\mathcal{P} = \begin{bmatrix} \mathcal{P}_{11} & \cdots & \mathcal{P}_{1n} \\ \vdots & \ddots & \vdots \\ \mathcal{P}_{m1} & \cdots & \mathcal{P}_{mn} \end{bmatrix}.$$

This means we can characterize a Markov process $(S_t)_t$ by a pair $(\mathcal{S}, \mathcal{P})$, where \mathcal{S} is the (finite) set of possible states and \mathcal{P} is the state transition matrix. For this reason, we shall also call $(\mathcal{S}, \mathcal{P})$ a **Markov process**.

Def'n 2.2. **Markov Reward Process**

A **Markov reward process** (MRP) is a tuple $(\mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma)$ such that

- (a) $(\mathcal{S}, \mathcal{P})$ is a Markov process;
- (b) \mathcal{R} is a **reward function** with

$$\mathcal{R}_s = \mathbb{E}(R_{t+1} | S_t = s)$$

for all $s \in \mathcal{S}$;¹ and

- (c) $\gamma \in [0, 1]$, called the **discount factor**.

¹ R_t represents the **reward** at step t .

Def'n 2.3. **Return**

The **return** from step t , denoted as G_t , is the total discounted reward from t :

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots.$$

According to the definition of G_t , the value of receiving reward R after $k+1$ steps is $\gamma^k R$, so G_t values immediate reward above delayed reward. Specifically, $\gamma \approx 0$ leads to *myopic* evaluation whereas $\gamma \approx 1$ leads to *far-sighted* evaluation.

One reason that we use a discount factor is because of the *uncertainty* in the future. To put this in another way, we *do not* have a perfect model of the environment.

Another reason is for mathematical convenience: by using discount rewards, we can avoid infinite returns (e.g. in cyclic Markov processes).

Def'n 2.4. **State Value Function**

The *state value function*, denoted as $v(s)$, of an MRP is the expectation of G_t starting from state s ;

$$v(s) = \mathbb{E}(G_t | S_t = s).$$